

GCP ACE Practise Set 07

Total points 0/0

Qs 200 onwards

Email *

tambetejashree02@gmail.com

✗ During a recent audit of your existing Google Cloud resources, you discovered several users with email addresses outside of your Google Workspace domain. You want to ensure that your resources are only shared with users whose email addresses match your domain. You need to remove any mismatched users, and you want to avoid having to audit your resources to identify mismatched users. What should you do? *

- A. Create a Cloud Scheduler task to regularly scan your projects and delete mismatched users.
- B. Create a Cloud Scheduler task to regularly scan your resources and delete mismatched users. ✗
- C. Set an organizational policy constraint to limit identities by domain to automatically remove mismatched users.
- D. Set an organizational policy constraint to limit identities by domain, and then retroactively remove the existing mismatched users

Correct answer

- D. Set an organizational policy constraint to limit identities by domain, and then retroactively remove the existing mismatched users



✓ All development (dev) teams in your organization are located in the United States. Each dev team has its own Google Cloud project. You want to restrict access so that each dev team can only create cloud resources in the United States (US). What should you do?

- A. Create a folder to contain all the dev projects. Create an organization policy ✓ to limit resources in US locations.
- B. Create an organization to contain all the dev projects. Create an Identity and Access Management (IAM) policy to limit the resources in US regions.
- C. Create an Identity and Access Management (IAM) policy to restrict the resources locations in the US. Apply the policy to all dev projects.
- D. Create an Identity and Access Management (IAM) policy to restrict the resources locations in all dev projects. Apply the policy to all dev roles.



✗ You have two subnets (subnet-a and subnet-b) in the default VPC. Your database servers are running in subnet-a. Your application servers and web servers are running in subnet-b. You want to configure a firewall rule that only allows database traffic from the application servers to the database servers. What should you do? *

- A. • Create service accounts sa-app and sa-db.
- Associate service account sa-app with the application servers and the service account sa-db with the database servers. ✓
- Create an ingress firewall rule to allow network traffic from source service account sa-app to target service account sa-db.
- B. • Create network tags app-server and db-server.
 - Add the app-server tag to the application servers and the db-server tag to the database servers.
 - Create an egress firewall rule to allow network traffic from source network tag app-server to target network tag db-server.
- C. • Create a service account sa-app and a network tag db-server.
 - Associate the service account sa-app with the application servers and the network tag db-server with the database servers.
 - Create an ingress firewall rule to allow network traffic from source VPC IP addresses and target the subnet-a IP addresses.
- D. • Create a network tag app-server and service account sa-db.
 - Add the tag to the application servers and associate the service account with the database servers.
 - Create an egress firewall rule to allow network traffic from source network tag app-server to target service account sa-db.

Correct answer

- A. • Create service accounts sa-app and sa-db.
- Associate service account sa-app with the application servers and the service account sa-db with the database servers.



- Create an ingress firewall rule to allow network traffic from source service account sa-app to target service account sa-db.

✓ You are building a data lake on Google Cloud for your Internet of Things (IoT) application. The IoT application has millions of sensors that are constantly streaming structured and unstructured data to your backend in the cloud. You want to build a highly available and resilient architecture based on Google-recommended practices. What should you do? *

- A. Stream data to Pub/Sub, and use Dataflow to send data to Cloud Storage. ✓
- B. Stream data to Pub/Sub, and use Storage Transfer Service to send data to BigQuery.
- C. Stream data to Dataflow, and use Dataprep by Trifacta to send data to Bigtable.
- D. Stream data to Dataflow, and use Storage Transfer Service to send data to BigQuery.

✓ you are configuring Cloud DNS. You want to create DNS records to point home.mydomain.com, mydomain.com, and www.mydomain.com to the IP address of your Google Cloud load balancer. What should you do? *

- A. Create one CNAME record to point [mydomain.com](#) to the load balancer, and create two A records to point WWW and HOME to [mydomain.com](#) respectively.
- B. Create one CNAME record to point [mydomain.com](#) to the load balancer, and create two AAAA records to point WWW and HOME to [mydomain.com](#) respectively.
- C. Create one A record to point [mydomain.com](#) to the load balancer, and create two CNAME records to point WWW and HOME to [mydomain.com](#) respectively. ✓
- D. Create one A record to point [mydomain.com](#) to the load balancer, and create two NS records to point WWW and HOME to [mydomain.com](#) respectively.



✖ Your company completed the acquisition of a startup and is now merging the IT systems of both companies. The startup had a production Google Cloud project in their organization. You need to move this project into your organization and ensure that the project is billed to your organization. You want to accomplish this task with minimal effort. What should you do? *

- A. Use the projects.move method to move the project to your organization. Update the billing account of the project to that of your organization.
- B. Ensure that you have an Organization Administrator Identity and Access Management (IAM) role assigned to you in both organizations. Navigate to the Resource Manager in the startup's Google Cloud organization, and drag the project to your company's organization.
- C. Create a Private Catalog for the Google Cloud Marketplace, and upload the resources of the startup's production project to the Catalog. Share the Catalog with your organization, and deploy the resources in your company's project. ✗
- D. Create an infrastructure-as-code template for all resources in the project by using Terraform, and deploy that template to a new project in your organization. Delete the project from the startup's Google Cloud organization.

Correct answer

- A. Use the projects.move method to move the project to your organization. Update the billing account of the project to that of your organization.



✓ Your company is using Google Workspace to manage employee accounts. *
Anticipated growth will increase the number of personnel from 100 employees to 1,000 employees within 2 years. Most employees will need access to your company's Google Cloud account. The systems and processes will need to support 10x growth without performance degradation, unnecessary complexity, or security issues. What should you do?

- A. Migrate the users to Active Directory. Connect the Human Resources system to Active Directory. Turn on Google Cloud Directory Sync (GCDS) for Cloud Identity. Turn on Identity Federation from Cloud Identity to Active Directory.
- B. Organize the users in Cloud Identity into groups. Enforce multi-factor authentication in Cloud Identity. ✓
- C. Turn on identity federation between Cloud Identity and Google Workspace. Enforce multi-factor authentication for domain wide delegation.
- D. Use a third-party identity provider service through federation. Synchronize the users from Google Workplace to the third-party provider in real time.



You want to host your video encoding software on Compute Engine. Your user base is growing rapidly, and users need to be able to encode their videos at any time without interruption or CPU limitations. You must ensure that your encoding solution is highly available, and you want to follow Google-recommended practices to automate operations. What should you do?

- A. Deploy your solution on multiple standalone Compute Engine instances, and increase the number of existing instances when CPU utilization on Cloud Monitoring reaches a certain threshold.
- B. Deploy your solution on multiple standalone Compute Engine instances, and replace existing instances with high-CPU instances when CPU utilization on Cloud Monitoring reaches a certain threshold.
- C. Deploy your solution to an instance group, and increase the number of available instances whenever you see high CPU utilization in Cloud Monitoring.
- D. Deploy your solution to an instance group, and set the autoscaling based on CPU utilization.

Correct answer

- D. Deploy your solution to an instance group, and set the autoscaling based on CPU utilization.



✗ You are working for a startup that was officially registered as a business 6 * months ago. As your customer base grows, your use of Google Cloud increases. You want to allow all engineers to create new projects without asking them for their credit card information. What should you do?

- A. Create a Billing account, associate a payment method with it, and provide all project creators with permission to associate that billing account with their projects.
- B. Grant all engineers permission to create their own billing accounts for each new project.
- C. Apply for monthly invoiced billing, and have a single invoice for the project paid by the finance team. ✗
- D. Create a billing account, associate it with a monthly purchase order (PO), and send the PO to Google Cloud.

Correct answer

- A. Create a Billing account, associate a payment method with it, and provide all project creators with permission to associate that billing account with their projects.



✖ Your company runs one batch process in an on-premises server that takes * around 30 hours to complete. The task runs monthly, can be performed offline, and must be restarted if interrupted. You want to migrate this workload to the cloud while minimizing cost. What should you do?

- A. Create an Instance Template with Spot VMs On. Create a Managed Instance Group from the template and adjust Target CPU Utilization. Migrate the workload.
- B. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.
- C. Migrate the workload to a Google Kubernetes Engine cluster with Spot nodes. ✖
- D. Migrate the workload to a Compute Engine Spot VM.

Correct answer

- B. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.



✖ Your continuous integration and delivery (CI/CD) server can't execute Google Cloud actions in a specific project because of permission issues. You need to validate whether the used service account has the appropriate roles in the specific project. What should you do? *

- A. Open the Google Cloud console, and check the Identity and Access Management (IAM) roles assigned to the service account at the project or inherited from the folder or organization levels.
- B. Open the Google Cloud console, and check the organization policies. ✖
- C. Open the Google Cloud console, and run a query to determine which resources this service account can access.
- D. Open the Google Cloud console, and run a query of the audit logs to find permission denied errors for this service account.

Correct answer

- A. Open the Google Cloud console, and check the Identity and Access Management (IAM) roles assigned to the service account at the project or inherited from the folder or organization levels.



✖ You recently discovered that your developers are using many service account keys during their development process. While you work on a long term improvement, you need to quickly implement a process to enforce short-lived service account credentials in your company. You have the following requirements:

- All service accounts that require a key should be created in a centralized project called pj-sa.
- Service account keys should only be valid for one day.

You need a Google-recommended solution that minimizes cost. What should you do?

- A. Implement a Cloud Run job to rotate all service account keys periodically in pj-sa.
- B. Implement a Kubernetes CronJob to rotate all service account keys periodically. Disable attachment of service accounts to resources in all projects with an exception to pj-sa. ✖
- C. Enforce an org policy constraint allowing the lifetime of service account keys to be 24 hours. Enforce an org policy constraint denying service account key creation with an exception on pj-sa.
- D. Enforce a DENY org policy constraint over the lifetime of service account keys for 24 hours. Disable attachment of service accounts to resources in all projects with an exception to pj-sa

Correct answer

- C. Enforce an org policy constraint allowing the lifetime of service account keys to be 24 hours. Enforce an org policy constraint denying service account key creation with an exception on pj-sa.



✓ Your application is running on Google Cloud in a managed instance group (MIG). You see errors in Cloud Logging for one VM that one of the processes is not responsive. You want to replace this VM in the MIG quickly. What should you do? *

- A. Use the gcloud compute instances update command with a REFRESH action for the VM.
- B. Use the gcloud compute instance-groups managed recreate-instances command to recreate the VM. ✓
- C. Select the MIG from the Compute Engine console and, in the menu, select Replace VMs.
- D. Update and apply the instance template of the MIG.

✗ You are running out of primary internal IP addresses in a subnet for a custom mode VPC. The subnet has the IP range 10.0.0.0/20, and the IP addresses are primarily used by virtual machines in the project. You need to provide more IP addresses for the virtual machines. What should you do? *

- A. Add a secondary IP range 10.1.0.0/20 to the subnet. ✗
- B. Change the subnet IP range from 10.0.0.0/20 to 10.0.0.0/18.
- C. Change the subnet IP range from 10.0.0.0/20 to 10.0.0.0/22.
- D. Convert the subnet IP range from IPv4 to IPv6.

Correct answer

- B. Change the subnet IP range from 10.0.0.0/20 to 10.0.0.0/18.



✖ Your company requires all developers to have the same permissions, regardless of the Google Cloud project they are working on. Your company's security policy also restricts developer permissions to Compute Engine, Cloud Functions, and Cloud SQL. You want to implement the security policy with minimal effort. What should you do? *

- A. • Create a custom role with Compute Engine, Cloud Functions, and Cloud SQL permissions in one project within the Google Cloud organization. • Copy the role across all projects created within the organization with the gcloud iam roles copy command. • Assign the role to developers in those projects.
- B. • Add all developers to a Google group in Google Groups for Workspace. • Assign the predefined role of Compute Admin to the Google group at the Google Cloud organization level.
- C. • Add all developers to a Google group in Cloud Identity. • Assign predefined roles for Compute Engine, Cloud Functions, and Cloud SQL permissions to the Google group for each project in the Google Cloud organization. ✗
- D. • Add all developers to a Google group in Cloud Identity. • Create a custom role with Compute Engine, Cloud Functions, and Cloud SQL permissions at the Google Cloud organization level. • Assign the custom role to the Google group.

Correct answer

- D. • Add all developers to a Google group in Cloud Identity. • Create a custom role with Compute Engine, Cloud Functions, and Cloud SQL permissions at the Google Cloud organization level. • Assign the custom role to the Google group.



✓ Your team is running an on-premises ecommerce application. The application contains a complex set of microservices written in Python, and each microservice is running on Docker containers. Configurations are injected by using environment variables. You need to deploy your current application to a serverless Google Cloud cloud solution. What should you do? *

- A. Use your existing CI/CD pipeline. Use the generated Docker images and deploy them to Cloud Run. Update the configurations and the required endpoints. ✓
- B. Use your existing continuous integration and delivery (CI/CD) pipeline. Use the generated Docker images and deploy them to Cloud Function. Use the same configuration as on-premises.
- C. Use the existing codebase and deploy each service as a separate Cloud Function. Update the configurations and the required endpoints.
- D. Use your existing codebase and deploy each service as a separate Cloud Run. Use the same configurations as on-premises.

✓ Your company has multiple projects linked to a single billing account in Google Cloud. You need to visualize the costs with specific metrics that should be dynamically calculated based on company-specific criteria. You want to automate the process. What should you do? *

- A. In the Google Cloud console, visualize the costs related to the projects in the Reports section.
- B. In the Google Cloud console, visualize the costs related to the projects in the Cost breakdown section.
- C. In the Google Cloud console, use the export functionality of the Cost table. Create a Looker Studio dashboard on top of the CSV export.
- D. Configure Cloud Billing data export to BigQuery for the billing account. Create a Looker Studio dashboard on top of the BigQuery export. ✓

✗ **The DevOps group in your organization needs full control of Compute Engine resources in your development project. However, they should not have permission to create or update any other resources in the project. You want to follow Google's recommendations for setting permissions for the DevOps group. What should you do?** *

- A. Grant the basic role roles/viewer and the predefined role roles/compute.admin to the DevOps group.
- B. Create an IAM policy and grant all compute.instanceAdmin.* permissions to the policy. Attach the policy to the DevOps group. ✗
- C. Create a custom role at the folder level and grant all compute.instanceAdmin.* permissions to the role. Grant the custom role to the DevOps group.
- D. Grant the basic role roles/editor to the DevOps group

Correct answer

- A. Grant the basic role roles/viewer and the predefined role roles/compute.admin to the DevOps group.

✗ **You have deployed an application on a single Compute Engine instance. The application writes logs to disk. Users start reporting errors with the application. You want to diagnose the problem. What should you do?** *

- A. Navigate to Cloud Logging and view the application logs.
- B. Configure a health check on the instance and set a “consecutive successes” Healthy threshold value of 1.
- C. Connect to the instance’s serial console and read the application logs. ✗
- D. Install and configure the Ops agent and view the logs from Cloud Logging.

Correct answer

- D. Install and configure the Ops agent and view the logs from Cloud Logging.

- ✖ You want to set up a Google Kubernetes Engine cluster. Verifiable node identity and integrity are required for the cluster, and nodes cannot be accessed from the internet. You want to reduce the operational cost of managing your cluster, and you want to follow Google-recommended practices. What should you do? *
- A. Deploy a private autopilot cluster.
 - B. Deploy a public autopilot cluster.
 - C. Deploy a standard public cluster and enable shielded nodes.
 - D. Deploy a standard private cluster and enable shielded nodes

Correct answer

- A. Deploy a private autopilot cluster.



✖ Your company wants to migrate their on-premises workloads to Google Cloud.

The current on-premises workloads consist of:

- A Flask web application
- A backend API
- A scheduled long-running background job for ETL and reporting

You need to keep operational costs low. You want to follow Google-recommended practices to migrate these workloads to serverless solutions on Google Cloud. What should you do

- A. Migrate the web application to App Engine and the backend API to Cloud Run.
Use Cloud Tasks to run your background job on Compute Engine.
- B. Migrate the web application to App Engine and the backend API to Cloud Run.
Use Cloud Tasks to run your background job on Cloud Run.
- C. Run the web application on a Cloud Storage bucket and the backend API on ✖
Cloud Run. Use Cloud Tasks to run your background job on Cloud Run.
- D. Run the web application on a Cloud Storage bucket and the backend API on
Cloud Run. Use Cloud Tasks to run your background job on Compute Engine.

Correct answer

- B. Migrate the web application to App Engine and the backend API to Cloud Run.
Use Cloud Tasks to run your background job on Cloud Run.



✗ You have created an application that is packaged into a Docker image. You * want to deploy the Docker image as a workload on Google Kubernetes Engine. What should you do?

- A. Upload the image to Cloud Storage and create a Kubernetes Service referencing the image.
- B. Upload the image to Cloud Storage and create a Kubernetes Deployment ✗ referencing the image.
- C. Upload the image to Artifact Registry and create a Kubernetes Service referencing the image.
- D. Upload the image to Artifact Registry and create a Kubernetes Deployment referencing the image.

Correct answer

- D. Upload the image to Artifact Registry and create a Kubernetes Deployment referencing the image.



✖ You are working in a team that has developed a new application that needs * to be deployed on Kubernetes. The production application is business critical and should be optimized for reliability. You need to provision a Kubernetes cluster and want to follow Google-recommended practices. What should you do?

- A. Create a GKE Autopilot cluster. Enroll the cluster in the rapid release channel.
- B. Create a GKE Autopilot cluster. Enroll the cluster in the stable release channel.
- C. Create a zonal GKE standard cluster. Enroll the cluster in the stable release channel. ✖
- D. Create a regional GKE standard cluster. Enroll the cluster in the rapid release channel.

Correct answer

- B. Create a GKE Autopilot cluster. Enroll the cluster in the stable release channel.



✖ Your company is running a three-tier web application on virtual machines that use a MySQL database. You need to create an estimated total cost of cloud infrastructure to run this application on Google Cloud instances and Cloud SQL. What should you do? *

- A. Create a Google spreadsheet with multiple Google Cloud resource combinations.
- B. On a separate sheet, import the current Google Cloud prices and use these prices for the calculations within formulas.
- C. Use the Google Cloud Pricing Calculator and select the Cloud Operations template to define your web application with as much detail as possible. ✖
- D. Implement a similar architecture on Google Cloud, and run a reasonable load test on a smaller scale. Check the billing information, and calculate the estimated costs based on the real load your system usually handles.
- E. Use the Google Cloud Pricing Calculator to determine the cost of every Google Cloud resource you expect to use. Use similar size instances for the web server, and use your current on-premises machines as a comparison for Cloud SQL.

Correct answer

- D. Use the Google Cloud Pricing Calculator to determine the cost of every Google Cloud resource you expect to use. Use similar size instances for the web server, and use your current on-premises machines as a comparison for Cloud SQL.



✗ **An external member of your team needs list access to compute images and * disks in one of your projects. You want to follow Google-recommended practices when you grant the required permissions to this user. What should you do?**

- A. Create a custom role, and add all the required compute.disks.list and compute.images.list permissions as includedPermissions. Grant the custom role to the user at the project level.
- B. Create a custom role based on the Compute Image User role. Add the compute.disks.list to the includedPermissions field. Grant the custom role to the user at the project level. ✗
- C. Create a custom role based on the Compute Storage Admin role. Exclude unnecessary permissions from the custom role. Grant the custom role to the user at the project level.
- D. Grant the Compute Storage Admin role at the project level.

Correct answer

- A. Create a custom role, and add all the required compute.disks.list and compute.images.list permissions as includedPermissions. Grant the custom role to the user at the project level.

✓ **You want to permanently delete a Pub/Sub topic managed by Config Connector in your Google Cloud project. What should you do? ***

- A. Use kubectl to create the label deleted-by-cnrm and to change its value to true for the topic resource.
- B. Use kubectl to delete the topic resource. ✓
- C. Use gcloud CLI to delete the topic.
- D. Use gcloud CLI to update the topic label managed-by-cnrm to false.



✗ You are migrating a business critical application from your local data center into Google Cloud. As part of your high-availability strategy, you want to ensure that any data used by the application will be immediately available if a zonal failure occurs. What should you do? *

- A. Store the application data on a zonal persistent disk. Create a snapshot schedule for the disk. If an outage occurs, create a new disk from the most recent snapshot and attach it to a new VM in another zone.
- B. Store the application data on a zonal persistent disk. If an outage occurs, create an instance in another zone with this disk attached.
- C. Store the application data on a regional persistent disk. Create a snapshot schedule for the disk. If an outage occurs, create a new disk from the most recent snapshot and attach it to a new VM in another zone. ✗
- D. Store the application data on a regional persistent disk. If an outage occurs, create an instance in another zone with this disk attached.

Correct answer

- D. Store the application data on a regional persistent disk. If an outage occurs, create an instance in another zone with this disk attached.



✗ You recently received a new Google Cloud project with an attached billing account where you will work. You need to create instances, set firewalls, and store data in Cloud Storage. You want to follow Google-recommended practices. What should you do? *

- A. Use the gcloud CLI services enable cloudresourcemanager.googleapis.com command to enable all resources. ✗
- B. Use the gcloud services enable compute.googleapis.com command to enable Compute Engine and the gcloud services enable storage-api.googleapis.com command to enable the Cloud Storage APIs.
- C. Open the Google Cloud console and enable all Google Cloud APIs from the API dashboard.
- D. Open the Google Cloud console and run gcloud init --project in a Cloud Shell.

Correct answer

- B. Use the gcloud services enable compute.googleapis.com command to enable Compute Engine and the gcloud services enable storage-api.googleapis.com command to enable the Cloud Storage APIs.

✓ Your team is using Linux instances on Google Cloud. You need to ensure that your team logs in to these instances in the most secure and cost efficient way. What should you do? *

- A. Attach a public IP to the instances and allow incoming connections from the internet on port 22 for SSH.
- B. Use the gcloud compute ssh command with the --tunnel-through-iap flag. Allow ingress traffic from the IP range 35.235.240.0/20 on port 22. ✓
- C. Use a third party tool to provide remote access to the instances.
- D. Create a bastion host with public internet access. Create the SSH tunnel to the instance through the bastion host.



✓ You are working for a hospital that stores its medical images in an on-premises data room. The hospital wants to use Cloud Storage for archival storage of these images. The hospital wants an automated process to upload any new medical images to Cloud Storage. You need to design and implement a solution. What should you do? *

- A. Create a Pub/Sub topic, and enable a Cloud Storage trigger for the Pub/Sub topic. Create an application that sends all medical images to the Pub/Sub topic.
- B. Create a script that uses the gcloud storage command to synchronize the on-premises storage with Cloud Storage, Schedule the script as a cron job. ✓
- C. Create a Pub/Sub topic, and create a Cloud Function connected to the topic that writes data to Cloud Storage. Create an application that sends all medical images to the Pub/Sub topic.
- D. In the Google Cloud console, go to Cloud Storage. Upload the relevant images to the appropriate bucket.

✓ Your company has a large quantity of unstructured data in different file formats. You want to perform ETL transformations on the data. You need to make the data accessible on Google Cloud so it can be processed by a Dataflow job. What should you do? *

- A. Upload the data to BigQuery using the bq command line tool.
- B. Upload the data to Cloud Storage using the gcloud storage command. ✓
- C. Upload the data into Cloud SQL using the import function in the Google Cloud console.
- D. Upload the data into Cloud Spanner using the import function in the Google Cloud console.



✓ Your application development team has created Docker images for an application that will be deployed on Google Cloud. Your team does not want to manage the infrastructure associated with this application. You need to ensure that the application can scale automatically as it gains popularity. What should you do? *

- A. Create an instance template with the container image, and deploy a Managed Instance Group with Autoscaling.
- B. Upload Docker images to Artifact Registry, and deploy the application on Google Kubernetes Engine using Standard mode.
- C. Upload Docker images to the Cloud Storage, and deploy the application on Google Kubernetes Engine using Standard mode.
- D. Upload Docker images to Artifact Registry, and deploy the application on Cloud Run. ✓



✖ You are responsible for a web application on Compute Engine. You want your support team to be notified automatically if users experience high latency for at least 5 minutes. You need a Google-recommended solution with no development cost. What should you do? *

- A. Export Cloud Monitoring metrics to BigQuery and use a Looker Studio dashboard to monitor your web application's latency.
- B. Create an alert policy to send a notification when the HTTP response latency exceeds the specified threshold.
- C. Implement an App Engine service which invokes the Cloud Monitoring API and sends a notification in case of anomalies.
- D. Use the Cloud Monitoring dashboard to observe latency and take the necessary actions when the response latency exceeds the specified threshold. ✖

Correct answer

- B. Create an alert policy to send a notification when the HTTP response latency exceeds the specified threshold.



✖ Your team wants to deploy a specific content management system (CMS) solution to Google Cloud. You need a quick and easy way to deploy and install the solution. What should you do? *

- A. Search for the CMS solution in Google Cloud Marketplace. Use gcloud CLI to deploy the solution.
- B. Search for the CMS solution in Google Cloud Marketplace. Deploy the solution directly from Cloud Marketplace.
- C. Search for the CMS solution in Google Cloud Marketplace. Use Terraform and the Cloud Marketplace ID to deploy the solution with the appropriate parameters.
- D. Use the installation guide of the CMS provider. Perform the installation through your configuration management system. ✗

Correct answer

- B. Search for the CMS solution in Google Cloud Marketplace. Deploy the solution directly from Cloud Marketplace.



- ✗ You used the gcloud container clusters command to create two Google Cloud Kubernetes (GKE) clusters: prod-cluster and dev-cluster.
- prod-cluster is a standard cluster.
 - dev-cluster is an auto-pilot cluster.

When you run the kubectl get nodes command, you only see the nodes from prod-cluster. Which commands should you run to check the node status for dev-cluster?

- A. gcloud container clusters get-credentials dev-cluster
- kubectl get nodes
- B. gcloud container clusters update -generate-password dev-cluster kubectl get nodes X
- C. kubectl config set-context dev-cluster
- kubectl cluster-info
- D. kubectl config set-credentials dev-cluster
- kubectl cluster-info
- Option 8

Correct answers

- A. gcloud container clusters get-credentials dev-cluster
- kubectl get nodes



- ✓ You are planning to migrate the following on-premises data management solutions to Google Cloud: *

- One MySQL cluster for your main database
- Apache Kafka for your event streaming platform
- One Cloud SQL for PostgreSQL database for your analytical and reporting needs

You want to implement Google-recommended solutions for the migration. You need to ensure that the new solutions provide global scalability and require minimal operational and infrastructure management. What should you do?

- A. Migrate from MySQL to Cloud SQL, from Kafka to Pub/Sub, and from Cloud SQL for PostgreSQL to BigQuery.
- B. Migrate from MySQL to Cloud Spanner, from Kafka to Pub/Sub, and from Cloud SQL for PostgreSQL to BigQuery. ✓
- C. Migrate from MySQL to Cloud Spanner, from Kafka to Memorystore, and from Cloud SQL for PostgreSQL to Cloud SQL.
- D. Migrate from MySQL to Cloud SQL, from Kafka to Memorystore, and from Cloud SQL for PostgreSQL to Cloud SQL.



✓ You have an on-premises data analytics set of binaries that processes data * files in memory for about 45 minutes every midnight. The sizes of those data files range from 1 gigabyte to 16 gigabytes. You want to migrate this application to Google Cloud with minimal effort and cost. What should you do?

- A. Create a container for the set of binaries. Use Cloud Scheduler to start a Cloud Run job for the container.
- B. Create a container for the set of binaries. Deploy the container to Google Kubernetes Engine (GKE) and use the Kubernetes scheduler to start the application.
- C. Upload the code to Cloud Functions. Use Cloud Scheduler to start the application.
- D. Lift and shift to a VM on Compute Engine. Use an instance schedule to start ✓ and stop the instance.

✓ You are running a web application on Cloud Run for a few hundred users. * Some of your users complain that the initial web page of the application takes much longer to load than the following pages. You want to follow Google's recommendations to mitigate the issue. What should you do?

- A. Set the minimum number of instances for your Cloud Run service to 3. ✓
- B. Set the concurrency number to 1 for your Cloud Run service.
- C. Set the maximum number of instances for your Cloud Run service to 100.
- D. Update your web application to use the protocol HTTP/2 instead of HTTP/1.1.



✓ You have a batch workload that runs every night and uses a large number of virtual machines (VMs). It is fault-tolerant and can tolerate some of the VMs being terminated. The current cost of VMs is too high. What should you do? *

- A. Run a test using simulated maintenance events. If the test is successful, use ✓ Spot N2 Standard VMs when running future jobs.
- B. Run a test using simulated maintenance events. If the test is successful, use N2 Standard VMs when running future jobs.
- C. Run a test using a managed instance group. If the test is successful, use N2 Standard VMs in the managed instance group when running future jobs.
- D. Run a test using N1 standard VMs instead of N2. If the test is successful, use N1 Standard VMs when running future jobs.



✖ You have a Bigtable instance that consists of three nodes that store personally identifiable information (PII) data. You need to log all read or write operations, including any metadata or configuration reads of this database table, in your company's Security Information and Event Management (SIEM) system. What should you do? *

- A. • Navigate to Cloud Monitoring in the Google Cloud console, and create a custom monitoring job for the Bigtable instance to track all changes. • Create an alert by using webhook endpoints, with the SIEM endpoint as a receiver.
- B. • Navigate to the Audit Logs page in the Google Cloud console, and enable Admin Write logs for the Bigtable instance. ✖
- C. • Navigate to the Audit Logs page in the Google Cloud console, and enable Data Read, Data Write and Admin Read logs for the Bigtable instance. • Create a Pub/Sub topic as a Cloud Logging sink destination, and add your SIEM as a subscriber to the topic.
- D. • Install the Ops Agent on the Bigtable instance during configuration. • Create a service account with read permissions for the Bigtable instance. • Create a custom Dataflow job with this service account to export logs to the company's SIEM system

Correct answer

- C. • Navigate to the Audit Logs page in the Google Cloud console, and enable Data Read, Data Write and Admin Read logs for the Bigtable instance. • Create a Pub/Sub topic as a Cloud Logging sink destination, and add your SIEM as a subscriber to the topic.



- ✗ You have an application that runs on Compute Engine VM instances in a custom Virtual Private Cloud (VPC). Your company's security policies only allow the use of internal IP addresses on VM instances and do not let VM instances connect to the internet. You need to ensure that the application can access a file hosted in a Cloud Storage bucket within your project.
- What should you do?

- A. Enable Private Service Access on the Cloud Storage Bucket.
- B. Add storage.googleapis.com to the list of restricted services in a VPC Service Controls perimeter and add your project to the list of protected projects. ✗
- C. Enable Private Google Access on the subnet within the custom VPC.
- D. Deploy a Cloud NAT instance and route the traffic to the dedicated IP address of the Cloud Storage bucket.

Correct answer

- C. Enable Private Google Access on the subnet within the custom VPC.

This content is neither created nor endorsed by Google. - [Terms of Service](#) - [Privacy Policy](#)

Google Forms





